



## TP1 : Apache Hadoop

*Rapport*

Réalisé par :

Abderrahim ALAKOUCHE  
Abdellah AGHLALOU

Encadré par :

Mme. Dounia ZAIDOUNI

# PLAN

Introduction

I

Installation et configuration d'un nœud unique

II

Exécution du programme MapReduce « Word Count » dans le cluster à nœud unique

III

Installation et configuration d'un cluster multi-nœuds

IV

Exécution du programme MapReduce « Word Count » dans le cluster multi-nœuds

V

Extension : benchmark détaillé



Hadoop est une infrastructure logicielle à code source libre pour le stockage et le traitement à grande échelle d'ensembles de données dans un environnement informatique distribué. Il est sponsorisé par Apache Software Foundation. Il est conçu pour passer de serveurs uniques à des milliers de machines, chacune offrant des fonctions de calcul et de stockage locales.

Hadoop a été créé par Doug Cutting et Mike Cafarella en 2005. Cutting, qui travaillait chez Yahoo! à l'époque, l'a nommé après l'éléphant de jouet de son fils. Hadoop a été initialement développé pour prendre en charge la distribution du projet de moteur de recherche.<sup>[1]</sup>

### Les objectifs de ce TP sont les suivants :

- Installation et configuration d'un nœud unique d'Apache Hadoop 3.2.1
- Exécution du programme MapReduce « **Word Count** » dans le cluster à nœud unique de Hadoop
- Installation et configuration d'un cluster multi-nœuds d'Apache Hadoop 3.2.1
- Exécution du programme MapReduce « **Word Count** » dans le cluster multi-nœuds de Hadoop
- Faire un benchmark sur les différentes plateformes cloud permettant la création d'un cluster multi-nœuds déjà configuré avec les outils de l'écosystème Hadoop.

[1]: <https://riptutorial.com/fr/hadoop>

Etape1 : Création d'un utilisateur hduser

```
alakouche@alakouche-VB:~$ sudo adduser hduser  
[sudo] password for alakouche:  
Adding user `hduser' ...  
Adding new group `hduser' (1001) ...  
Adding new user `hduser' (1001) with group `hduser' ...  
Creating home directory `/home/hduser' ...  
Copying files from `/etc/skel' ...  
Enter new UNIX password:  
Retype new UNIX password:  
passwd: password updated successfully  
Changing the user information for hduser  
Enter the new value, or press ENTER for the default  
  Full Name []: Alakouche/Aghlalou  
  Room Number []:  
  Work Phone []:  
  Home Phone []:  
  Other []:  
Is the information correct? [Y/n] Y  
alakouche@alakouche-VB:~$ sudo adduser hduser sudo  
Adding user `hduser' to group `sudo' ...  
Adding user hduser to group sudo  
Done.  
alakouche@alakouche-VB:~$
```

Création d'un compte normal (non root) **hduser**

Cette commande va nous permettre d'éviter les erreurs du types « **hduser is not in the sudoers file** »

```
hduser@alakouche-VB:~$
```

On redémarre la machine virtuelle afin de basculer vers le compte **hduser**

## Etape2 : Mise en place de la clé ssh

```
hduser@alakouche-VB:~$ sudo apt-get install openssh-server  
[sudo] password for hduser:  
Reading package lists... Done  
Building dependency tree  
Reading state information... Done  
The following packages were automatically installed and are no longer required:  
  efibootmgr libeal1-mesa libfwud1 liblvm9 libwvland-eal1-mesa
```

Installation de paquet nécessaire pour ssh.

```
hduser@alakouche-VB:~$ ssh-keygen -t rsa -P ""  
Generating public/private rsa key pair.  
Enter file in which to save the key (/home/hduser/.ssh/id_rsa):  
Created directory '/home/hduser/.ssh'.  
Your identification has been saved in /home/hduser/.ssh/id_rsa.  
Your public key has been saved in /home/hduser/.ssh/id_rsa.pub.  
The key fingerprint is:  
SHA256:01DBfrKXsTqdvPCMWg+XQx0xLEKtb5M8K33UP6qWIIfA hduser@alakouche-VB  
The key's randomart image is:  
+---[RSA 2048]---+  
|   o+ .o |  
|   .o o .o |  
|   .. o .. |  
|   . o+ o.. . |  
|   S .*.=.. |  
|   E.o@... . |  
|   =B** . |  
|   .+@*.... |  
|   ..o+*+. . |  
+---[SHA256]---+  
hduser@alakouche-VB:~$ cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys  
hduser@alakouche-VB:~$ chmod 0600 ~/.ssh/authorized_keys
```

Mettre en place la clé ssh pour son propre compte.

Hadoop nécessite un accès SSH pour gérer les différents nœuds. Bien que nous soyons dans une configuration simple nœud, nous avons besoin de configurer l'accès vers localhost pour l'utilisateur hduser que nous venons de créer précédemment.

Autoriser l'accès au SSH de la machine avec cette nouvelle clé fraîchement créée

```
hduser@alakouche-VB:~$ ssh-copy-id -i /home/hduser/.ssh/id_rsa.pub hduser@localhost
/usr/bin/ssh-copy-id: INFO: Source of key(s) to be installed: "/home/hduser/.ssh/id_rsa.pub"
The authenticity of host 'localhost (127.0.0.1)' can't be established.
ECDSA key fingerprint is SHA256:QKauX3NvcSv1gtSlfSwVKABjuLCFVn/mfGTDqbndRbs.
Are you sure you want to continue connecting (yes/no)? yes
/usr/bin/ssh-copy-id: INFO: attempting to log in with the new key(s), to filter out any tha
/usr/bin/ssh-copy-id: WARNING: All keys were skipped because they already exist on the remo
(if you think this is a mistake, you may want to use -f option)
```

On copie la clé public sur le serveur localhost

```
hduser@alakouche-VB:~$ ssh localhost
Welcome to Ubuntu 18.04.5 LTS (GNU/Linux 5.4.0-53-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:     https://landscape.canonical.com
 * Support:        https://ubuntu.com/advantage

 * Canonical Livepatch is available for installation.
   - Reduce system reboots and improve kernel security. Activate at:
     https://ubuntu.com/livepatch

0 packages can be updated.
0 updates are security updates.

New release '20.04.1 LTS' available.
Run 'do-release-upgrade' to upgrade to it.

Your Hardware Enablement Stack (HWE) is supported until April 2023.
Last login: Mon Nov 16 20:06:02 2020 from 127.0.0.1
hduser@alakouche-VB:~$ exit
logout
Connection to localhost closed.
```

On teste la connexion à localhost

### Etape 3 : Installation de JAVA 8

- Hadoop nécessite l'installation de Java. Pour ce TP, la version 8 de Java sera utilisée via la distribution **OpenJDK**

```
hduser@alakouche-VB:~$  
hduser@alakouche-VB:~$ su -  
Password:  
root@alakouche-VB:~# mkdir /opt/java
```

Création de répertoire **/opt/java**

```
root@alakouche-VB:~# cd /home/hduser/Documents  
root@alakouche-VB:/home/hduser/Documents# ls  
jdk-8u71-linux-x64.tar.gz  
root@alakouche-VB:/home/hduser/Documents# tar -zvxf jdk-8u71-linux-x64.tar.gz  
jdk1.8.0_71/  
jdk1.8.0_71/db/  
jdk1.8.0_71/db/lib/  
jdk1.8.0_71/db/lib/derbyLocale_pl.jar
```

On va ensuite décompresser l'archive  
**jdk8u71linuxx64.tar.gz**

```
root@alakouche-VB:/home/hduser/Documents# ls  
jdk1.8.0_71 jdk-8u71-linux-x64.tar.gz  
root@alakouche-VB:/home/hduser/Documents# mv jdk1.8.0_71/ /opt/java/
```

On déplace le jdk vers **/opt/java/**

On informe le système où java et ses exécutables sont installés.

```
root@alakouche-VB:/home/hduser/Documents# cd /opt/java/jdk1.8.0_71/
root@alakouche-VB:/opt/java/jdk1.8.0_71# update-alternatives --install /usr/bin/java java /opt/java/jdk1.8.0_71/bin/java 100
root@alakouche-VB:/opt/java/jdk1.8.0_71# update-alternatives --config java
There is only one alternative in link group java (providing /usr/bin/java): /opt/java/jdk1.8.0_71/bin/java
Nothing to configure.
root@alakouche-VB:/opt/java/jdk1.8.0_71#
root@alakouche-VB:/opt/java/jdk1.8.0_71# update-alternatives --install /usr/bin/javac javac /opt/java/jdk1.8.0_71/bin/javac 100
update-alternatives: using /opt/java/jdk1.8.0_71/bin/javac to provide /usr/bin/javac (javac) in auto mode
root@alakouche-VB:/opt/java/jdk1.8.0_71# update-alternatives --config javac
There is only one alternative in link group javac (providing /usr/bin/javac): /opt/java/jdk1.8.0_71/bin/javac
Nothing to configure.
root@alakouche-VB:/opt/java/jdk1.8.0_71#
```

```
root@alakouche-VB:/opt/java/jdk1.8.0_71# apt install vim
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following packages were automatically installed and are no longer required:
  efibootmgr libegl1-mesa libfwupd1 libllvm9 libwayland-egl1-mesa
Use 'apt autoremove' to remove them.
The following additional packages will be installed:
  vim-runtime
```

Installation de l'éditeur de texte  
« Vim ».



```
root@alakouche-VB:/opt/java/jdk1.8.0_71# vim /etc/profile  
root@alakouche-VB:/opt/java/jdk1.8.0_71#
```

```
done  
unset i  
fi  
export JAVA_HOME=/opt/java/jdk1.8.0_71/  
export JRE_HOME=/opt/java/jdk1.8.0_71/jre  
export PATH=$PATH:/opt/java/jdk1.8.0_71/bin:/opt/java/jdk1.8.0_71/jre/bin  
-- INSERT --
```

Mettre en place de manière permanente les variables d'environnement JAVA pour tous les utilisateurs

On recharge le fichier **/etc/profile**

```
root@alakouche-VB:~# source /etc/profile  
root@alakouche-VB:~# su - hduser  
hduser@alakouche-VB:~$ source /etc/profile  
hduser@alakouche-VB:~$ echo $PATH  
/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/bin:/usr/games:/usr/local/games:/snap/bin:/opt/java/jdk1.8.0_71/bin:/opt/java/jdk1.8.0_71/jre/bin:/opt/java/jdk1.8.0_71/jre/bin:  
hduser@alakouche-VB:~$
```

On teste la mise en place des variables d'environnement dans le terminal hadoop

Assurent que la version Java est correctement installée

```
hduser@alakouche-VB:~$ java -version  
java version "1.8.0_71"  
Java(TM) SE Runtime Environment (build 1.8.0_71-b15)  
Java HotSpot(TM) 64-Bit Server VM (build 25.71-b15, mixed mode)
```

#### Etape 4 : Installation d'Apache Hadoop 3.2.1

```
hduser@alakouche-VB:~/Documents$ tar -zxvf hadoop-3.2.1.tar.gz
```

On décomprime l'archive : hadoop3.1.2.tar.gz

```
hduser@alakouche-VB:~/Documents$ ls  
hadoop-3.2.1  hadoop-3.2.1.tar.gz  jdk-8u71-linux-x64.tar.gz  
hduser@alakouche-VB:~/Documents$ mv hadoop-3.2.1 hadoop  
hduser@alakouche-VB:~/Documents$ sudo mv hadoop /usr/local/hadoop/  
[sudo] password for hduser:
```

```
hduser@alakouche-VB:~/Documents$ sudo chown -R hduser /usr/local/hadoop  
hduser@alakouche-VB:~/Documents$
```

On affecte les droits à notre utilisateur **hduse**.

```
hduser@alakouche-VB:~/Documents$ sudo mkdir -p /usr/local/hadoop_store/hdfs/namenode  
hduser@alakouche-VB:~/Documents$ sudo mkdir -p /usr/local/hadoop_store/hdfs/datanode  
hduser@alakouche-VB:~/Documents$ sudo chown -R hduser /usr/local/hadoop_store  
hduser@alakouche-VB:~/Documents$
```

Création de datanode et namenode.  
On affecte ensuite les droits à notre utilisateur **hduse**.

## Etape 5 : Configuration d'Apache Hadoop 3.2.1

- Il faut maintenant définir la configuration de Hadoop et pour cela plusieurs fichiers de configurations doivent être modifiés

```
hduser@alakouche-VB:~/Documents$ vim .bashrc  
hduser@alakouche-VB:~/Documents$
```

```
fi  
#HADOOP VARIABLES START  
export JAVA_HOME=/opt/java/jdk1.8.0_71/  
export HADOOP_INSTALL=/usr/local/hadoop  
export PATH=$PATH:$HADOOP_INSTALL/bin  
export PATH=$PATH:$HADOOP_INSTALL/sbin  
export HADOOP_MAPRED_HOME=$HADOOP_INSTALL  
export HADOOP_COMMON_HOME=$HADOOP_INSTALL  
export HADOOP_HDFS_HOME=$HADOOP_INSTALL  
export YARN_HOME=$HADOOP_INSTALL  
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_INSTALL/lib/native  
#export HADOOP_OPTS="-Djava.library.path=$HADOOP_INSTALL/lib"  
#HADOOP VARIABLES END  
.x
```

```
hduser@alakouche-VB:~$ cd /usr/local/hadoop/etc/hadoop  
hduser@alakouche-VB:/usr/local/hadoop/etc/hadoop$ vim hadoop-env.sh
```

```
# JAVA_HOME=/usr/java/testing hdfs dfs -ls  
#  
# export JAVA_HOME=/opt/java/jdk1.8.0_71/ # Therefore, the vast ma  
# are configured for substitution and not append. If append  
# is preferable, modify this file accordingly.  
  
###
```

```
hduser@alakouche-VB:~$ sudo mkdir -p /app/hadoop/tmp  
hduser@alakouche-VB:~$ sudo chown hduser /app/hadoop/tmp  
hduser@alakouche-VB:~$
```

Création de répertoire des fichiers temporaires de Hadoop

- Modification des autres fichiers de configurations : on ajoute les lignes suivantes entre les balises de configurations.

```
<configuration>
  <property>
    <name>hadoop.tmp.dir</name>
    <value>/app/hadoop/tmp</value>
  </property>
  <property>
    <name>fs.default.name</name>
    <value>hdfs://localhost:54310</value>
  </property>
</configuration>
"core-site.xml" 28L, 958C
```

```
<configuration>
  <property>
    <name>dfs.replication</name>
    <value>1</value>
  </property>
  <property>
    <name>dfs.namenode.name.dir</name>
    <value>file:/usr/local/hadoop_store/hdfs/namenode</value>
  </property>
  <property>
    <name>dfs.datanode.data.dir</name>
    <value>file:/usr/local/hadoop_store/hdfs/datanode</value>
  </property>
</configuration>
```

```
hduser@alakouche-VB:~$ cd /usr/local/hadoop/etc/hadoop/
hduser@alakouche-VB:/usr/local/hadoop/etc/hadoop$ vim core-site.xml
hduser@alakouche-VB:/usr/local/hadoop/etc/hadoop$ vim hdfs-site.xml
hduser@alakouche-VB:/usr/local/hadoop/etc/hadoop$ vim mapred-site.xml
hduser@alakouche-VB:/usr/local/hadoop/etc/hadoop$ vim yarn-site.xml
```

```
<configuration>
  <!-- Site specific YARN configuration properties -->
  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
  </property>
</configuration>
```

```
<!-- Put site-specific property overrides -->

<configuration>
  <property>
    <name>mapred.job.tracker</name>
    <value>localhost:54311</value>
  </property>
</configuration>
```

- Avant de démarrer le serveur Hadoop, il faut formater le système de fichiers HDFS.

```
hduser@alakouche-VB:/usr/local/hadoop/etc/hadoop$ hdfs namenode -format
2020-11-18 10:30:03,876 INFO namenode.NameNode: STARTUP_MSG:
/*****
STARTUP_MSG: Starting NameNode
STARTUP_MSG:   host = alakouche-VB/127.0.1.1
```

```
2020-11-18 10:30:10,322 INFO util.GSet: 0.029999999329447746% max memory 953.2 MB = 292.8 KB
2020-11-18 10:30:10,325 INFO util.GSet: capacity      = 2^15 = 32768 entries
2020-11-18 10:30:10,553 INFO namenode.FSImage: Allocated new BlockPoolId: BP-1419653612-127.0.1.1-1605695410488
2020-11-18 10:30:10,685 INFO common.Storage: Storage directory /usr/local/hadoop_store/hdfs/namenode has been successfully formatted.
2020-11-18 10:30:10,944 INFO namenode.FSImageFormatProtobuf: Saving image file /usr/local/hadoop_store/hdfs/namenode/current/fsimage.c
0000000000000000 using no compression
2020-11-18 10:30:11,550 INFO namenode.FSImageFormatProtobuf: Image file /usr/local/hadoop_store/hdfs/namenode/current/fsimage.ckpt_000
00000000 of size 398 bytes saved in 0 seconds .
2020-11-18 10:30:11,631 INFO namenode.NNStorageRetentionManager: Going to retain 1 images with txid >= 0
2020-11-18 10:30:11,653 INFO namenode.FSImage: FSImageSaver clean checkpoint: txid=0 when meet shutdown.
2020-11-18 10:30:11,655 INFO namenode.NameNode: SHUTDOWN_MSG:
/*****
SHUTDOWN_MSG: Shutting down NameNode at alakouche-VB/127.0.1.1
*****/
hduser@alakouche-VB:/usr/local/hadoop/etc/hadoop$
```

- Démarrage de Hadoop

```
hduser@alakouche-VB:/usr/local/hadoop/etc/hadoop$ start-dfs.sh
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [alakouche-VB]
alakouche-VB: Warning: Permanently added 'alakouche-vb' (ECDSA) to the list of known hosts.
2020-11-18 10:42:44,628 WARN util.NativeCodeLoader: Unable to load native-hadoop library
applicable
```

```
hduser@alakouche-VB:/usr/local/hadoop/etc/hadoop$ start-yarn.sh
Starting resourcemanager
Starting nodemanagers
hduser@alakouche-VB:/usr/local/hadoop/etc/hadoop$ █
```

```
hduser@alakouche-VB:/usr/local/hadoop/etc/hadoop$ jps
4773 DataNode
5447 NodeManager
4986 SecondaryNameNode
4619 NameNode
5819 Jps
5295 ResourceManager
```

- Vérification de l'installation  
On peut maintenant accéder à l'interface web Hadoop .

- interface graphique accessible par notre navigateur

java - "hdfs: command no" X    Namenode information X    +

localhost:9870/dfshealth.html#tab-overview 80% ... ☰ ⚡

Hadoop Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities ▾

## Overview 'localhost:54310' (active)

<b>Started:</b>	Wed Nov 18 11:42:28 +0100 2020
<b>Version:</b>	3.2.1, rb3cbbb467e22ea829b3808f4b7b01d07e0bf3842
<b>Compiled:</b>	Tue Sep 10 16:56:00 +0100 2019 by rohitsharmaks from branch-3.2.1
<b>Cluster ID:</b>	CID-6c81b3d0-31ec-4b4f-ad67-d53e850f5863
<b>Block Pool ID:</b>	BP-1419653612-127.0.1.1-1605695410488

## Summary

Security is off.

Safemode is off.

1 files and directories, 0 blocks (0 replicated blocks, 0 erasure coded block groups) = 1 total filesystem object(s).

Heap Memory used 41.72 MB of 61.88 MB Heap Memory. Max Heap Memory is 953.19 MB.

Non Heap Memory used 45.67 MB of 46.98 MB Committed Non Heap Memory. Max Non Heap Memory is <unbounded>.

<b>Configured Capacity:</b>	9.78 GB
<b>Configured Remote Capacity:</b>	0 B

- également on peut visualiser l'avancement et les résultats de notre Jobs

A screenshot of a Linux desktop environment showing the Apache Hadoop Cluster UI. The window title is "All Applications". The URL in the address bar is "localhost:8088/cluster". The left sidebar has a "Cluster Metrics" section with tabs for "About", "Nodes", "Node Labels", "Applications" (which is selected), and "Scheduler". The "Applications" tab shows metrics like Apps Submitted (0), Apps Pending (0), Apps Running (0), Apps Completed (0), Containers Running (0), Memory Used (0 B), Memory Total (0 B), Memory Reserved (0 B), VCores Used (0), VCores Total (0), and VCores Reserved (0). Below this are sections for "Cluster Nodes Metrics" and "Scheduler Metrics". The Scheduler Metrics table shows the Capacity Scheduler with Scheduling Resource Type as [memory-mb (unit=Mi), vcores] and Minimum Allocation as <memory:1024, vCores:1>. The table header includes columns for ID, User, Name, Application Type, Queue, Application Priority, StartTime, LaunchTime, FinishTime, State, FinalStatus, Running Containers, Allocated CPU VCores, Allocated Memory MB, Reserved CPU VCores, Reserved Memory MB, % of Queue, % of Cluster, Progress, Tracking UI, and Blacklisted Nodes. A message "No data available in table" is displayed. At the bottom, it says "Showing 0 to 0 of 0 entries" and provides navigation links for First, Previous, Next, and Last. The desktop taskbar at the bottom shows icons for various applications like Firefox, File Explorer, and Microsoft Word.

```
hduser@alakouche-VB:~$ hdfs dfsadmin -report  
2020-11-18 13:05:37,920 WARN util.NativeCodeLoader: Unabi  
applicable
```

```
-----  
Live datanodes (1):  
  
Name: 127.0.0.1:9866 (localhost)  
Hostname: alakouche-VB  
Decommission Status : Normal  
Configured Capacity: 10499674112 (9.78 GB)  
DFS Used: 32768 (32 KB)  
Non DFS Used: 9009983488 (8.39 GB)  
DFS Remaining: 936116224 (892.75 MB)  
DFS Used%: 0.00%  
DFS Remaining%: 8.92%  
Configured Cache Capacity: 0 (0 B)  
Cache Used: 0 (0 B)  
Cache Remaining: 0 (0 B)  
Cache Used%: 100.00%  
Cache Remaining%: 0.00%  
Xceivers: 1  
Last contact: Wed Nov 18 13:05:41 WET 2020  
Last Block Report: Wed Nov 18 13:01:22 WET 2020  
Num of Blocks: 0  
  
hduser@alakouche-VB:~$
```

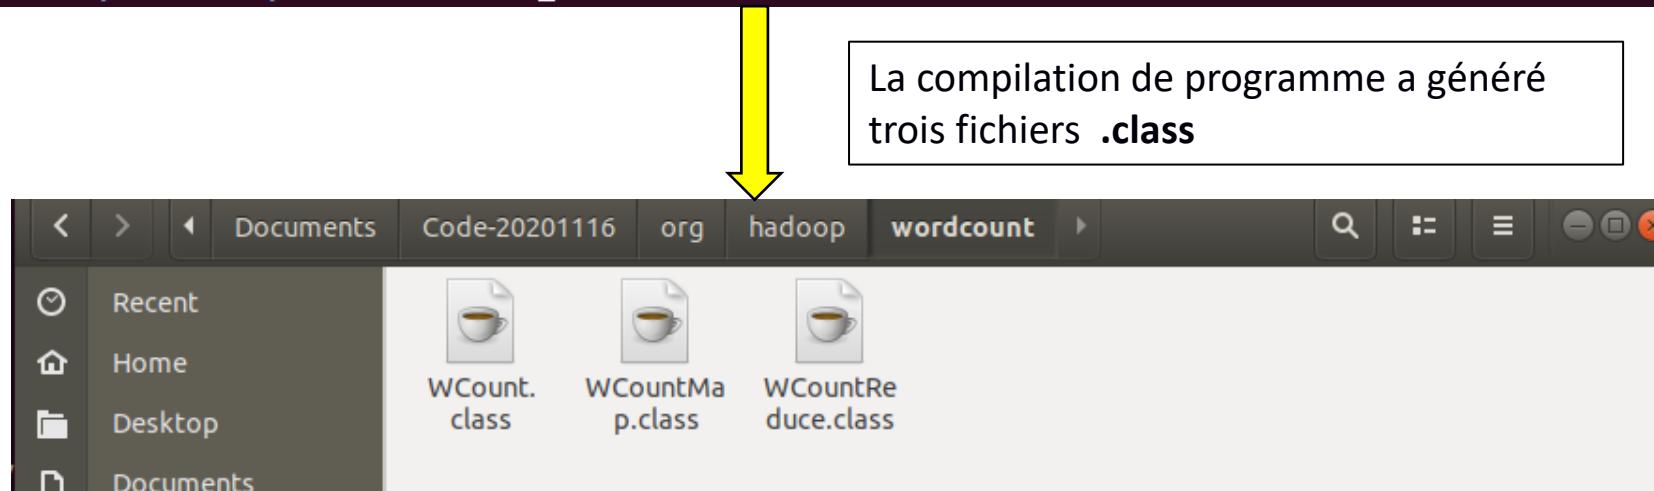
On lance cette commande afin de vérifier le bon fonctionnement de notre noeud

```
hduser@alakouche-VB:~$ cd /home/hduser/Documents/  
hduser@alakouche-VB:~/Documents$ ls  
Code-20201116  Code-20201116.zip  hadoop-3.2.1.tar.gz  jdk-8u71-linux-x64.tar.gz  poeme  'Script export classpath'  
hduser@alakouche-VB:~/Documents$ cd Code-20201116/  
hduser@alakouche-VB:~/Documents/Code-20201116$ mkdir -p org/hadoop/wordcount/  
hduser@alakouche-VB:~/Documents/Code-20201116$
```

```
hduser@alakouche-VB:~$ cd /home/hduser/Documents/  
hduser@alakouche-VB:~/Documents$ ls  
Code-20201116  Code-20201116.zip  hadoop-3.2.1.tar.gz  jdk-8u71-linux-x64.tar.gz  poeme  'Script export classpath'  
hduser@alakouche-VB:~/Documents$ cd Code-20201116/  
hduser@alakouche-VB:~/Documents/Code-20201116$ mkdir -p org/hadoop/wordcount/  
hduser@alakouche-VB:~/Documents/Code-20201116$ sudo chmod -R 777 org/hadoop/wordcount/  
[sudo] password for hduser:  
hduser@alakouche-VB:~/Documents/Code-20201116$ javac -classpath ${HADOOP_CLASSPATH} WCount*.java  
hduser@alakouche-VB:~/Documents/Code-20201116$ mv *.class org/hadoop/wordcount/  
hduser@alakouche-VB:~/Documents/Code-20201116$
```

Compilation de programme

La compilation de programme a généré  
trois fichiers .class



### On génère le jar file

```
hduser@alakouche-VB:~/Documents/Code-20201116$ jar -cvf wcount.jar . /home/hduser/Documents/code/org  
/home/hduser/Documents/code/org : no such file or directory  
added manifest  
adding: WCount.java(in = 2005) (out= 890)(deflated 55%)  
adding: WCountReduce.java(in = 1142) (out= 589)(deflated 48%)  
adding: WCountMap.java(in = 1057) (out= 565)(deflated 46%)  
adding: org/(in = 0) (out= 0)(stored 0%)  
adding: org/hadoop/(in = 0) (out= 0)(stored 0%)  
adding: org/hadoop/wordcount/(in = 0) (out= 0)(stored 0%)  
adding: org/hadoop/wordcount/WCountReduce.class(in = 1834) (out= 775)(deflated 57%)  
adding: org/hadoop/wordcount/WCount.class(in = 1646) (out= 859)(deflated 47%)  
adding: org/hadoop/wordcount/WCountMap.class(in = 1674) (out= 722)(deflated 56%)  
hduser@alakouche-VB:~/Documents/Code-20201116$
```

```
hduser@alakouche-VB:/usr/local/hadoop$ bin/hdfs dfs -put /home/hduser/Documents/code/poeme.txt /
2020-11-18 21:32:17,495 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
2020-11-18 21:32:21,446 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
hduser@alakouche-VB:/usr/local/hadoop$ bin/hdfs dfs -ls /
2020-11-18 21:32:58,571 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
-rw-r--r-- 1 hduser supergroup 1670 2020-11-18 21:32 /poeme.txt
hduser@alakouche-VB:/usr/local/hadoop$
```



Hadoop Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities ▾

## Browse Directory

/	Go!						
Show 25 entries	Search:						
Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	hduser	supergroup	1.63 KB	Nov 18 22:32	1	128 MB	poeme.txt

Showing 1 to 1 of 1 entries

Previous 1 Next

```
hduser@alakouche-VB:/usr/local/hadoop$ cd ~  
hduser@alakouche-VB:~$ mv /home/hduser/Documents/Code-20201116 /home/hduser/Documents/code  
hduser@alakouche-VB:~$ cd /home/hduser/Documents/code/  
hduser@alakouche-VB:~/Documents/code$ hadoop jar wcount.jar org.hadoop.wordcount.WCount /poeme.txt /results  
2020-11-18 21:37:20,705 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-jav  
applicable  
2020-11-18 21:37:23,809 INFO impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
```

On renomme le répertoire pour plus de simplicité

```
WRONG_LENGTH=0  
WRONG_MAP=0  
WRONG_REDUCE=0  
File Output Format Counters  
    Bytes Written=2823  
2020-11-18 21:37:32,129 INFO mapred.LocalJobRunner: Finishing task: attempt_local535396870_0001_r_000000_0  
2020-11-18 21:37:32,134 INFO mapred.LocalJobRunner: reduce task executes complete.  
2020-11-18 21:37:32,642 INFO mapreduce.Job: map 100% reduce 100%  
2020-11-18 21:37:32,643 INFO mapreduce.Job: Job job_local535396870_0001 completed successfully  
2020-11-18 21:37:32,702 INFO mapreduce.Job: Counters: 36  
File System Counters  
    FILE: Number of bytes read=28076  
    FILE: Number of bytes written=1072698  
    FILE: Number of read operations=0  
    FILE: Number of large read operations=0  
    FILE: Number of write operations=0  
    HDFS: Number of bytes read=3340  
    HDFS: Number of bytes written=2823  
    HDFS: Number of read operations=15  
    HDFS: Number of large read operations=0
```

On Vérifie la présence de fichier de résultats

```
hduser@alakouche-VB:~/Documents/code$ hadoop fs -ls /results  
2020-11-18 21:39:31,300 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where  
applicable  
Found 2 items  
-rw-r--r-- 1 hduser supergroup 0 2020-11-18 21:37 /results/_SUCCESS  
-rw-r--r-- 1 hduser supergroup 2823 2020-11-18 21:37 /results/part-r-00000
```

- Affichage de résultats

```
hduser@alakouche-VB:~/Documents/code$ hadoop fs -cat /results/part-r-00000
2020-11-18 21:40:22,517 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
2020-11-18 21:40:26,618 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
```

a 6 occurrences.  
adoraient 1 occurrences.  
ailes 1 occurrences.  
aima 1 occurrences.  
amour 1 occurrences.  
au 11 occurrences.  
bas 1 occurrences.  
belle 1 occurrences.  
bles 1 occurrences.  
bras 1 occurrences.  
bretagne 1 occurrences.  
brula 1 occurrences.  
celle 1 occurrences.  
celui 20 occurrences.  
cette 1 occurrences.  
chancelle 1 occurrences.  
chapelle 1 occurrences.  
ciel 10 occurrences.  
citadelle 1 occurrences.  
clarte 1 occurrences.  
coeur 2 occurrences.  
combat 1 occurrences.  
comment 1 occurrences.  
commun 1 occurrences.  
coule 2 occurrences.  
couleur 1 occurrences.  
court 1 occurrences.

meme 2 occurrences.  
mirabelle 1 occurrences.  
montait 1 occurrences.  
mourra 1 occurrences.  
muisse 1 occurrences.  
muscat 1 occurrences.  
ne 1 occurrences.  
nom 1 occurrences.  
nouvelle 1 occurrences.  
ny 10 occurrences.  
ou 3 occurrences.  
par 1 occurrences.  
pas 11 occurrences.  
passent 1 occurrences.  
plus 2 occurrences.  
pour 1 occurrences.  
prefere 1 occurrences.  
prison 1 occurrences.  
prisonniere 1 occurrences.  
qua 1 occurrences.  
quand 2 occurrences.  
quaucun 1 occurrences.  
que 2 occurrences.  
quelle 1 occurrences.  
querelles 1 occurrences.  
qui 25 occurrences.  
quil 1 occurrences.  
quimporte 1 occurrences.  
raisin 1 occurrences.  
rats 1 occurrences.  
rebelle 2 occurrences.  
rechantera 1 occurrences.  
repetant 1 occurrences.  
reseda 1 occurrences.

nom 1 occurrences.  
nouvelle 1 occurrences.  
ny 10 occurrences.  
ou 3 occurrences.  
par 1 occurrences.  
pas 11 occurrences.  
passent 1 occurrences.  
plus 2 occurrences.  
pour 1 occurrences.  
prefere 1 occurrences.  
prison 1 occurrences.  
prisonniere 1 occurrences.  
qua 1 occurrences.  
quand 2 occurrences.  
quaucun 1 occurrences.  
que 2 occurrences.  
quelle 1 occurrences.  
querelles 1 occurrences.  
qui 25 occurrences.  
quil 1 occurrences.  
quimporte 1 occurrences.  
raisin 1 occurrences.  
rats 1 occurrences.  
rebelle 2 occurrences.  
rechantera 1 occurrences.  
repetant 1 occurrences.  
reseda 1 occurrences.

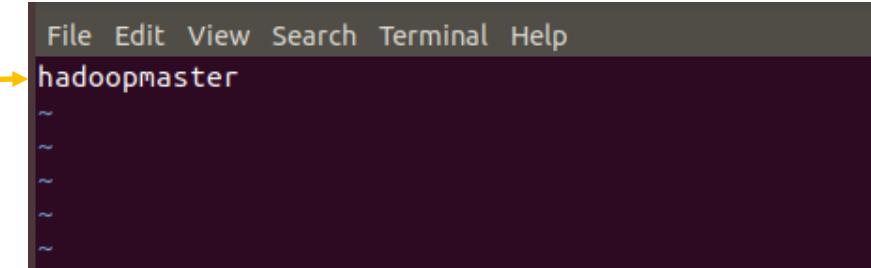
....

- On arrête tous les daemons en Cours d'exécution sur notre machine virtuelle.

```
hduser@alakouche-VB:~/Documents/code$ stop-dfs.sh
Stopping namenodes on [localhost]
Stopping datanodes
Stopping secondary namenodes [alakouche-VB]
2020-11-18 21:43:29,997 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java class
applicable
hduser@alakouche-VB:~/Documents/code$ stop-yarn.sh
Stopping nodemanagers
Stopping resourcemanager
hduser@alakouche-VB:~/Documents/code$
```

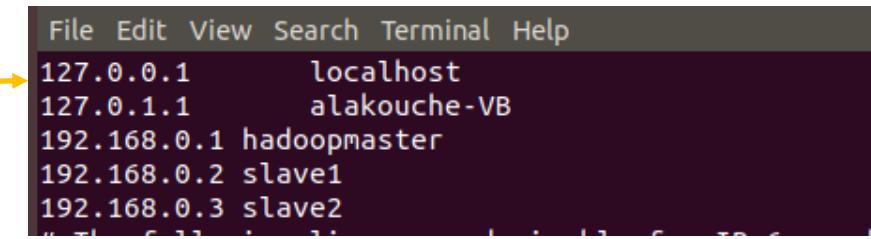
Nous allons travaillé avec la machine virtuelle précédemment configuré en node unique dans la section précédente.  
Et pour cela plusieurs fichiers doivent être modifiés

```
hduser@alakouche-VB:~$ sudo vim /etc/hostname  
[sudo] password for hduser:  
hduser@alakouche-VB:~$
```



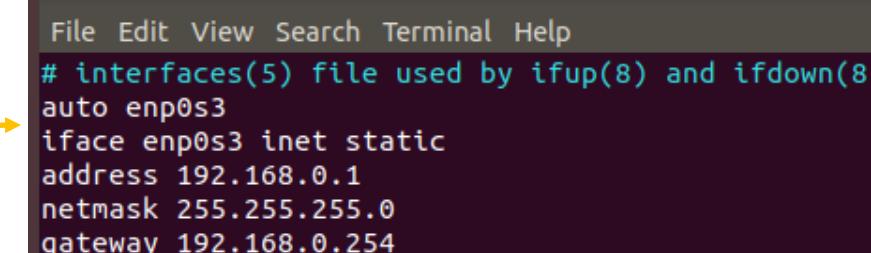
```
File Edit View Search Terminal Help  
hadoopmaster  
~  
~  
~  
~  
~
```

```
hduser@alakouche-VB:~$ sudo vim /etc/hosts  
hduser@alakouche-VB:~$
```



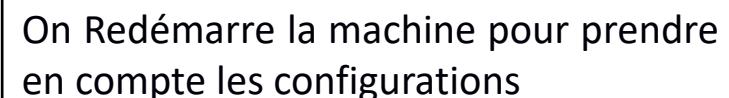
```
File Edit View Search Terminal Help  
127.0.0.1      localhost  
127.0.1.1      alakouche-VB  
192.168.0.1    hadoopmaster  
192.168.0.2    slave1  
192.168.0.3    slave2
```

```
hduser@alakouche-VB:~$ sudo vim /etc/network/interfaces  
hduser@alakouche-VB:~$
```



```
File Edit View Search Terminal Help  
# interfaces(5) file used by ifup(8) and ifdown(8)  
auto enp0s3  
iface enp0s3 inet static  
address 192.168.0.1  
netmask 255.255.255.0  
gateway 192.168.0.254
```

```
hduser@alakouche-VB:~$ sudo reboot
```



On Redémarre la machine pour prendre en compte les configurations

- Suppression des fichiers du répertoire de stockage de données créé par l'installation single node de Hadoop

```
hduser@hadoopmaster:~$ cd /usr/local/hadoop_store/
hduser@hadoopmaster:/usr/local/hadoop_store$ rm -rf *
hduser@hadoopmaster:/usr/local/hadoop_store$ mkdir -p /usr/local/hadoop_store/hdfs/namenode
hduser@hadoopmaster:/usr/local/hadoop_store$ chown -R hduser /usr/local/hadoop_store/hdfs/namenode
hduser@hadoopmaster:/usr/local/hadoop_store$
```

- Modification des fichiers de configuration de hadoop

```
hduser@hadoopmaster:/usr/local/hadoop_store$ cd /usr/local/hadoop/etc/hadoop/
hduser@hadoopmaster:/usr/local/hadoop/etc/hadoop$ sudo vim core-site.xml
hduser@hadoopmaster:/usr/local/hadoop/etc/hadoop$ sudo vim hdfs-site.xml
hduser@hadoopmaster:/usr/local/hadoop/etc/hadoop$ sudo vim mapred-site.xml
hduser@hadoopmaster:/usr/local/hadoop/etc/hadoop$ sudo vim yarn-site.xml
hduser@hadoopmaster:/usr/local/hadoop/etc/hadoop$
```

```
<configuration>
<property>
<name>yarn.resourcemanager.resource-tracker.address</name>
<value>hadoopmaster:8025</value>
</property>
<property>
<name>yarn.resourcemanager.scheduler.address</name>
<value>hadoopmaster:8030</value>
</property>
<property>
<name>yarn.resourcemanager.address</name>
<value>hadoopmaster:8050</value>
</property>
<property>
<name>yarn.nodemanager.aux-services</name>
<value>mapreduce_shuffle</value>
</property>
<property>
<name>yarn.nodemanager.aux-services.mapreduce.shuffle.class</name>
<value>org.apache.hadoop.mapred.ShuffleHandler</value>
</property>
</configuration>
```

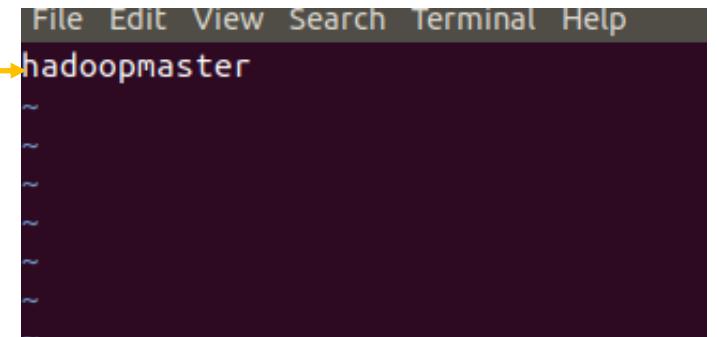
```
-->
<!-- Put site-specific property overrides in this file. -->
<configuration>
<property>
<name>hadoop.tmp.dir</name>
<value>/app/hadoop/tmp</value>
</property>
<property>
<name>fs.default.name</name>
<value>hdfs://hadoopmaster:54310</value>
</property>
</configuration>
~
```

```
<!-- Put site-specific property overrides in this file. -->
<configuration>
<property>
<name>dfs.replication</name>
<value>2</value>
</property>
<property>
<name>dfs.namenode.name.dir</name>
<value>file:/usr/local/hadoop_store/hdfs/namenode</value>
</property>
</configuration>
```

```
<!-- Put site-specific property overrides in this file. -->
<configuration>
<property>
<name>mapred.job.tracker</name>
<value>hadoopmaster:54311</value>
</property>
</configuration>
```

- On crée le fichier masters qui contient le hostname de la machine master

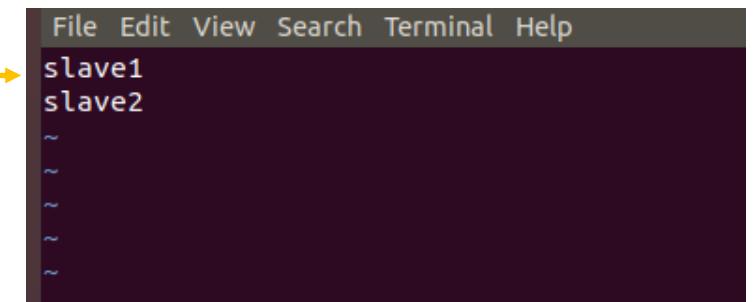
```
hduser@hadoopmaster:/usr/local/hadoop/etc/hadoop$ vim masters  
hduser@hadoopmaster:/usr/local/hadoop/etc/hadoop$
```



```
File Edit View Search Terminal Help  
hadoopmaster  
~  
~  
~  
~  
~  
~
```

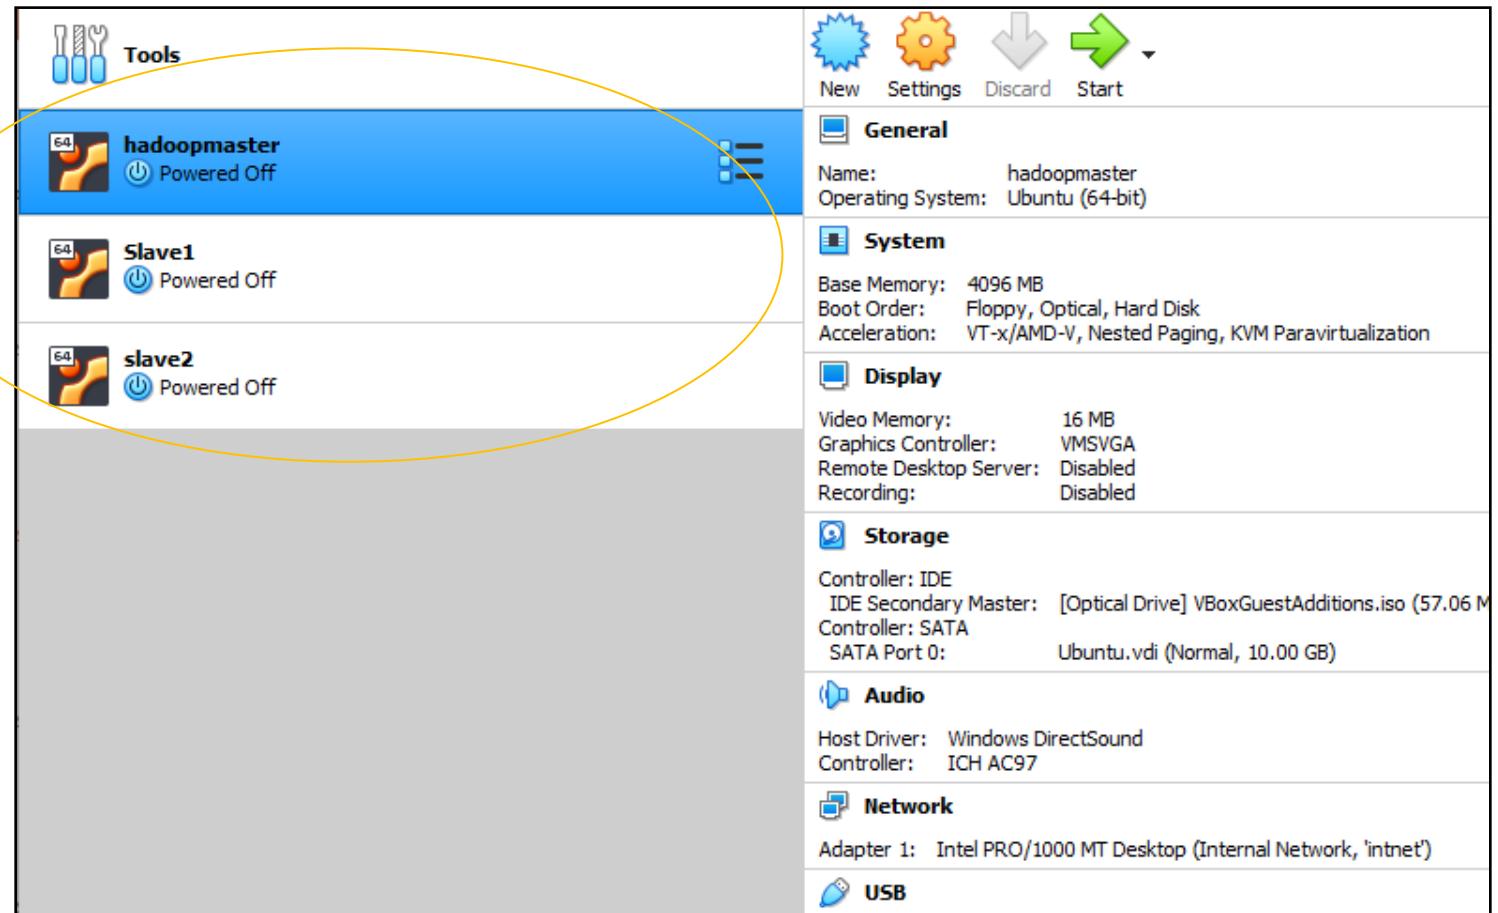
- Modification de fichier workers qui contient le hostname de chaque machine slave dans le répertoire

```
hduser@hadoopmaster:/usr/local/hadoop/etc/hadoop$ vim workers  
hduser@hadoopmaster:/usr/local/hadoop/etc/hadoop$
```



```
File Edit View Search Terminal Help  
slave1  
slave2  
~  
~  
~  
~  
~
```

## Clonage de la machine hadoopmaster





- Configuration de l'adresse IP fixe de la machine **slave1**

```
hduser@hadoopmaster:~$ sudo vim /etc/network/interfaces  
[sudo] password for hduser:  
hduser@hadoopmaster:~$
```

```
# interfaces(5) file used by ifup(8) and ifdown(8)  
auto enp0s3  
iface enp0s3 inet static  
    address 192.168.0.2  
    netmask 255.255.255.0  
    gateway 192.168.0.254  
~
```

- Modification de fichier /etc/hostname de la machine **slave1**

```
hduser@hadoopmaster:~$ sudo vim /etc/hostname  
hduser@hadoopmaster:~$
```

```
File Edit View Search Terminal Help  
slave1  
~  
~  
~  
~
```

```
hduser@slave1:~$
```

On Redémarre la machine pour prendre en compte les configurations.



- Configuration de l'adresse IP fixe de la machine **slave2**

```
hduser@hadoopmaster:~$ sudo vim /etc/network/interfaces  
[sudo] password for hduser:  
hduser@hadoopmaster:~$
```

```
File Edit View Search Terminal Help  
# interfaces(5) file used by ifup(8) and ifdown(8)  
auto enp0s3  
iface enp0s3 inet static  
address 192.168.0.3  
netmask 255.255.255.0  
gateway 192.168.0.254  
~
```

- Modification de fichier /etc/hostname de la machine **slave2**

```
hduser@hadoopmaster:~$ sudo vim /etc/hostname  
hduser@hadoopmaster:~$
```

```
File Edit View Search Terminal Help  
slave2  
~  
~  
~
```

```
hduser@slave2:~$
```

On Redémarre la machine pour prendre en compte les configurations.

- On Supprime les fichiers du répertoire de stockage de données créer par l'installation single node de Hadoop pour slave1 et slave2:

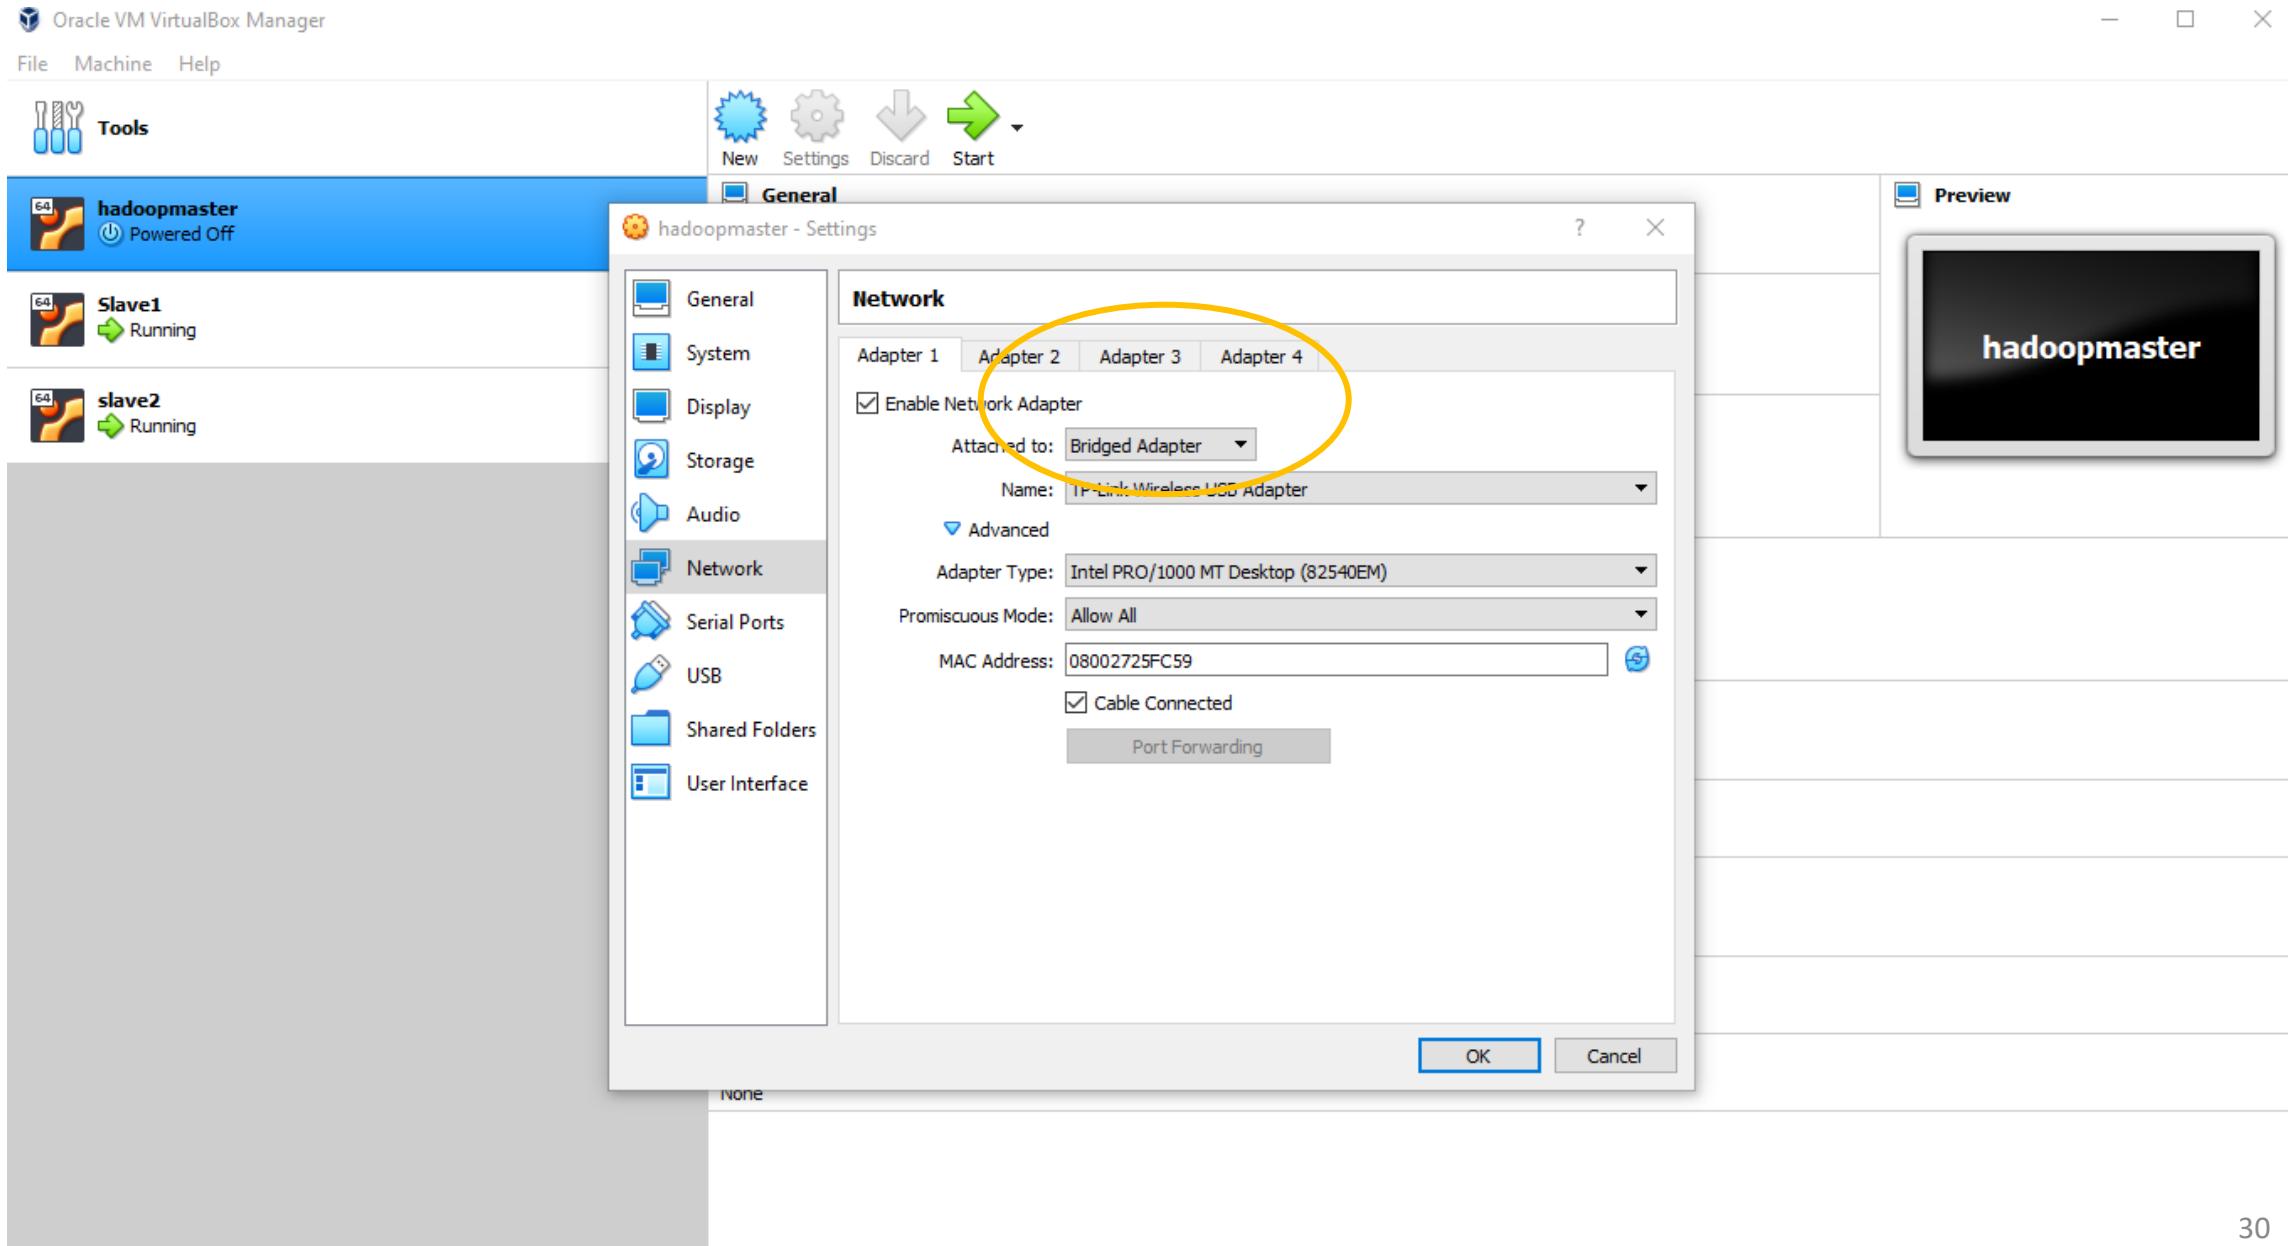
```
hduser@slave1:~$ cd /usr/local/hadoop_store/  
hduser@slave1:/usr/local/hadoop_store$ rm -rf *  
hduser@slave1:/usr/local/hadoop_store$ mkdir -p /usr/local/hadoop_store/hdfs/datanode  
hduser@slave1:/usr/local/hadoop_store$ chown -R hduser /usr/local/hadoop_store/hdfs/datanode  
hduser@slave1:/usr/local/hadoop_store$
```



```
hduser@slave2:~$ cd /usr/local/hadoop_store/  
hduser@slave2:/usr/local/hadoop_store$ rm -rf *  
hduser@slave2:/usr/local/hadoop_store$ mkdir -p /usr/local/hadoop_store/hdfs/datanode  
hduser@slave2:/usr/local/hadoop_store$  
hduser@slave2:/usr/local/hadoop_store$ chown -R hduser /usr/local/hadoop_store/hdfs/datanode  
hduser@slave2:/usr/local/hadoop_store$
```



- Connexion entre les machines du cluster



## Test de connexion entre slave1 et slave2

```
hduser@slave1:~$ ping 192.168.0.3
PING 192.168.0.3 (192.168.0.3) 56(84) bytes of data.
64 bytes from 192.168.0.3: icmp_seq=1 ttl=64 time=0.825 ms
64 bytes from 192.168.0.3: icmp_seq=2 ttl=64 time=1.17 ms
64 bytes from 192.168.0.3: icmp_seq=3 ttl=64 time=1.13 ms
64 bytes from 192.168.0.3: icmp_seq=4 ttl=64 time=1.20 ms
64 bytes from 192.168.0.3: icmp_seq=5 ttl=64 time=1.07 ms
64 bytes from 192.168.0.3: icmp_seq=6 ttl=64 time=1.08 ms
64 bytes from 192.168.0.3: icmp_seq=7 ttl=64 time=1.13 ms
64 bytes from 192.168.0.3: icmp_seq=8 ttl=64 time=1.04 ms
64 bytes from 192.168.0.3: icmp_seq=9 ttl=64 time=0.583 ms
64 bytes from 192.168.0.3: icmp_seq=10 ttl=64 time=1.10 ms
^C
--- 192.168.0.3 ping statistics ---
10 packets transmitted, 10 received, 0% packet loss, time 9017ms
rtt min/avg/max/mdev = 0.583/1.037/1.202/0.180 ms
hduser@slave1:~$
```

```
hduser@slave2:~$ ping 192.168.0.2
PING 192.168.0.2 (192.168.0.2) 56(84) bytes of data.
64 bytes from 192.168.0.2: icmp_seq=1 ttl=64 time=0.393 ms
64 bytes from 192.168.0.2: icmp_seq=2 ttl=64 time=1.12 ms
64 bytes from 192.168.0.2: icmp_seq=3 ttl=64 time=1.12 ms
64 bytes from 192.168.0.2: icmp_seq=4 ttl=64 time=0.994 ms
64 bytes from 192.168.0.2: icmp_seq=5 ttl=64 time=1.13 ms
64 bytes from 192.168.0.2: icmp_seq=6 ttl=64 time=1.09 ms
64 bytes from 192.168.0.2: icmp_seq=7 ttl=64 time=1.06 ms
64 bytes from 192.168.0.2: icmp_seq=8 ttl=64 time=1.26 ms
64 bytes from 192.168.0.2: icmp_seq=9 ttl=64 time=1.07 ms
^C
--- 192.168.0.2 ping statistics ---
9 packets transmitted, 9 received, 0% packet loss, time 8017ms
rtt min/avg/max/mdev = 0.393/1.030/1.261/0.235 ms
hduser@slave2:~$
```

## Test de connexion entre hadoopmaster et slave1

```
hduser@hadoopmaster:~$ ping 192.168.0.2
PING 192.168.0.2 (192.168.0.2) 56(84) bytes of data.
64 bytes from 192.168.0.2: icmp_seq=1 ttl=64 time=0.660 ms
64 bytes from 192.168.0.2: icmp_seq=2 ttl=64 time=1.02 ms
64 bytes from 192.168.0.2: icmp_seq=3 ttl=64 time=1.07 ms
64 bytes from 192.168.0.2: icmp_seq=4 ttl=64 time=0.930 ms
64 bytes from 192.168.0.2: icmp_seq=5 ttl=64 time=1.16 ms
64 bytes from 192.168.0.2: icmp_seq=6 ttl=64 time=1.02 ms
64 bytes from 192.168.0.2: icmp_seq=7 ttl=64 time=1.16 ms
^C
--- 192.168.0.2 ping statistics ---
7 packets transmitted, 7 received, 0% packet loss, time 6017ms
rtt min/avg/max/mdev = 0.660/1.006/1.166/0.160 ms
hduser@hadoopmaster:~$
```

```
hduser@slave1:~$ ping 192.168.0.1
PING 192.168.0.1 (192.168.0.1) 56(84) bytes of data.
64 bytes from 192.168.0.1: icmp_seq=1 ttl=64 time=0.773 ms
64 bytes from 192.168.0.1: icmp_seq=2 ttl=64 time=1.07 ms
64 bytes from 192.168.0.1: icmp_seq=3 ttl=64 time=1.05 ms
64 bytes from 192.168.0.1: icmp_seq=4 ttl=64 time=1.10 ms
64 bytes from 192.168.0.1: icmp_seq=5 ttl=64 time=1.17 ms
64 bytes from 192.168.0.1: icmp_seq=6 ttl=64 time=1.12 ms
^C
--- 192.168.0.1 ping statistics ---
6 packets transmitted, 6 received, 0% packet loss, time 5007ms
rtt min/avg/max/mdev = 0.773/1.052/1.173/0.132 ms
hduser@slave1:~$
```

## Test de connexion entre hadoopmaster et slave2

```
hduser@hadoopmaster:~$ ping 192.168.0.3
PING 192.168.0.3 (192.168.0.3) 56(84) bytes of data.
64 bytes from 192.168.0.3: icmp_seq=1 ttl=64 time=0.368 ms
64 bytes from 192.168.0.3: icmp_seq=2 ttl=64 time=1.05 ms
64 bytes from 192.168.0.3: icmp_seq=3 ttl=64 time=0.975 ms
64 bytes from 192.168.0.3: icmp_seq=4 ttl=64 time=0.428 ms
64 bytes from 192.168.0.3: icmp_seq=5 ttl=64 time=0.382 ms
64 bytes from 192.168.0.3: icmp_seq=6 ttl=64 time=0.329 ms
64 bytes from 192.168.0.3: icmp_seq=7 ttl=64 time=0.680 ms
^C
--- 192.168.0.3 ping statistics ---
7 packets transmitted, 7 received, 0% packet loss, time 6082ms
rtt min/avg/max/mdev = 0.329/0.602/1.056/0.283 ms
hduser@hadoopmaster:~$
```

```
hduser@slave2:~$ ping 192.168.0.1
PING 192.168.0.1 (192.168.0.1) 56(84) bytes of data.
64 bytes from 192.168.0.1: icmp_seq=1 ttl=64 time=7.16 ms
64 bytes from 192.168.0.1: icmp_seq=2 ttl=64 time=4.27 ms
64 bytes from 192.168.0.1: icmp_seq=3 ttl=64 time=3.26 ms
64 bytes from 192.168.0.1: icmp_seq=4 ttl=64 time=3.27 ms
64 bytes from 192.168.0.1: icmp_seq=5 ttl=64 time=2.55 ms
64 bytes from 192.168.0.1: icmp_seq=6 ttl=64 time=3.58 ms
64 bytes from 192.168.0.1: icmp_seq=7 ttl=64 time=11.7 ms
64 bytes from 192.168.0.1: icmp_seq=8 ttl=64 time=3.59 ms
64 bytes from 192.168.0.1: icmp_seq=9 ttl=64 time=3.03 ms
64 bytes from 192.168.0.1: icmp_seq=10 ttl=64 time=6.81 ms
^C
--- 192.168.0.1 ping statistics ---
10 packets transmitted, 10 received, 0% packet loss, time 9017ms
rtt min/avg/max/mdev = 2.558/4.931/11.757/2.719 ms
hduser@slave2:~$
```

- On copie la clé ssh pour configurer un accès ssh sans mot de passe entre les machines du cluster.

```
hduser@hadoopmaster:~$ ssh-copy-id -i /home/hduser/.ssh/id_rsa.pub hduser@hadoopmaster
/usr/bin/ssh-copy-id: INFO: Source of key(s) to be installed: "/home/hduser/.ssh/id_rsa.pub"
/usr/bin/ssh-copy-id: INFO: attempting to log in with the new key(s), to filter out any that are already installed
/usr/bin/ssh-copy-id: WARNING: All keys were skipped because they already exist on the remote system.
(if you think this is a mistake, you may want to use -f option)

hduser@hadoopmaster:~$
```

```
hduser@hadoopmaster:~$ ssh-copy-id -i /home/hduser/.ssh/id_rsa.pub hduser@slave1
/usr/bin/ssh-copy-id: INFO: Source of key(s) to be installed: "/home/hduser/.ssh/id_rsa.pub"
The authenticity of host 'slave1 (192.168.0.2)' can't be established.
ECDSA key fingerprint is SHA256:QKauX3NvcSv1gtSlfSwVKABjuLCFVn/mfGTDqbndRbs.
Are you sure you want to continue connecting (yes/no)? yes
/usr/bin/ssh-copy-id: INFO: attempting to log in with the new key(s), to filter out any that are already installed
/usr/bin/ssh-copy-id: WARNING: All keys were skipped because they already exist on the remote system.
(if you think this is a mistake, you may want to use -f option)

hduser@hadoopmaster:~$
```

```
hduser@hadoopmaster:~$ ssh-copy-id -i /home/hduser/.ssh/id_rsa.pub hduser@slave2
/usr/bin/ssh-copy-id: INFO: Source of key(s) to be installed: "/home/hduser/.ssh/id_rsa.pub"
The authenticity of host 'slave2 (192.168.0.3)' can't be established.
ECDSA key fingerprint is SHA256:QKauX3NvcSv1gtSlfSwVKABjuLCFVn/mfGTDqbndRbs.
Are you sure you want to continue connecting (yes/no)? yes
/usr/bin/ssh-copy-id: INFO: attempting to log in with the new key(s), to filter out any that are already installed
/usr/bin/ssh-copy-id: WARNING: All keys were skipped because they already exist on the remote system.
(if you think this is a mistake, you may want to use -f option)

hduser@hadoopmaster:~$
```

- On Teste la connexion **ssh** entre les machines du cluster.

```
hduser@hadoopmaster:~$ ssh slave1
Welcome to Ubuntu 18.04.5 LTS (GNU/Linux 5.4.0-53-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:     https://landscape.canonical.com
 * Support:        https://ubuntu.com/advantage

 * Canonical Livepatch is available for installation.
   - Reduce system reboots and improve kernel security. Activate at:
     https://ubuntu.com/livepatch

8 packages can be updated.
0 updates are security updates.

Failed to connect to https://changelogs.ubuntu.com/meta-release-lts. Ch

Your Hardware Enablement Stack (HWE) is supported until April 2023.
Last login: Mon Nov 16 21:14:22 2020 from 127.0.0.1
hduser@slave1:~$ █
```

```
hduser@hadoopmaster:~$ ssh slave2
Welcome to Ubuntu 18.04.5 LTS (GNU/Linux 5.4.0-53-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:     https://landscape.canonical.com
 * Support:        https://ubuntu.com/advantage

 * Canonical Livepatch is available for installation.
   - Reduce system reboots and improve kernel security. Activate at:
     https://ubuntu.com/livepatch

8 packages can be updated.
0 updates are security updates.

Failed to connect to https://changelogs.ubuntu.com/meta-release-lts. Ch

Your Hardware Enablement Stack (HWE) is supported until April 2023.
Last login: Mon Nov 16 21:14:22 2020 from 127.0.0.1
hduser@slave2:~$ █
```

- Modification de fichier **hdfs-site.xml** de la machine virtuelle slave1 et slave2

```
hduser@slave1:~$ cd /usr/local/hadoop/etc/hadoop  
hduser@slave1:/usr/local/hadoop/etc/hadoop$ sudo vim hdfs-site.xml  
hduser@slave1:/usr/local/hadoop/etc/hadoop$
```

```
hduser@slave2:~$ cd /usr/local/hadoop/etc/hadoop  
hduser@slave2:/usr/local/hadoop/etc/hadoop$ sudo vim hdfs-site.xml  
hduser@slave2:/usr/local/hadoop/etc/hadoop$
```

```
<!-- Put site-specific property overrides in this file. -->  
  
<configuration>  
  <property>  
    <name>dfs.replication</name>  
    <value>2</value>  
  </property>  
  <property>  
    <name>dfs.datanode.data.dir</name>  
    <value>file:/usr/local/hadoop_store/hdfs/datanode</value>  
  </property>  
</configuration>  
~  
~  
~
```

- Avant de démarrer le cluster, il faut vider aussi le répertoire de stockage du nœud hadoopmaster

```
hduser@hadoopmaster:~$ cd /usr/local/hadoop_store/  
hduser@hadoopmaster:/usr/local/hadoop_store$ rm -rf *
```

```
hduser@hadoopmaster:/usr/local/hadoop_store$ mkdir -p /usr/local/hadoop_store/hdfs/namenode  
hduser@hadoopmaster:/usr/local/hadoop_store$ chown -R hduser /usr/local/hadoop_store/  
hduser@hadoopmaster:/usr/local/hadoop_store$
```

- Avant de démarrer le serveur Hadoop, il faut formater le système de fichiers HDFS.

```
hduser@hadoopmaster:/usr/local/hadoop/etc/hadoop$ hdfs namenode -format  
2020-11-19 19:54:20,720 INFO namenode.NameNode: STARTUP_MSG:  
/*****
```

```
2020-11-18 22:45:46,078 INFO util.GSet: 0.029999999329447746% max memory 953.2 MB = 292.8 KB  
2020-11-18 22:45:46,078 INFO util.GSet: capacity      = 2^15 = 32768 entries  
2020-11-18 22:45:46,324 INFO namenode.FSImage: Allocated new BlockPoolId: BP-1427748328-192.168.0.1-1605739546281  
2020-11-18 22:45:46,511 INFO common.Storage: Storage directory /usr/local/hadoop_store/hdfs/namenode has been successfully formatted.  
2020-11-18 22:45:46,821 INFO namenode.FSImageFormatProtobuf: Saving image file /usr/local/hadoop_store/hdfs/namenode/current/fsimage.ck  
0000000000000000 using no compression  
2020-11-18 22:45:47,327 INFO namenode.FSImageFormatProtobuf: Image file /usr/local/hadoop_store/hdfs/namenode/current/fsimage.ckpt_0000  
00000000 of size 401 bytes saved in 0 seconds .  
2020-11-18 22:45:47,470 INFO namenode.NNStorageRetentionManager: Going to retain 1 images with txid >= 0  
2020-11-18 22:45:47,521 INFO namenode.FSImage: FSImageSaver clean checkpoint: txid=0 when meet shutdown.  
2020-11-18 22:45:47,522 INFO namenode.NameNode: SHUTDOWN_MSG:  
/*****  
SHUTDOWN_MSG: Shutting down NameNode at hadoopmaster/192.168.0.1  
*****/  
hduser@hadoopmaster:/usr/local/hadoop/etc/hadoop$
```

- Démarrage de Hadoop

```
hduser@hadoopmaster:/usr/local/hadoop/etc/hadoop$ start-dfs.sh
Starting namenodes on [hadoopmaster]
Starting datanodes
Starting secondary namenodes [hadoopmaster]
2020-11-19 20:42:23,107 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
hduser@hadoopmaster:/usr/local/hadoop/etc/hadoop$ start-yarn.sh
Starting resourcemanager
Starting nodemanagers
hduser@hadoopmaster:/usr/local/hadoop/etc/hadoop$
```

```
hduser@hadoopmaster:/usr/local/hadoop_store$ jps
18452 Jps
18166 ResourceManager
18327 NodeManager
17870 SecondaryNameNode
17663 DataNode
17503 NameNode
hduser@hadoopmaster:/usr/local/hadoop_store$
```

```
hduser@slave1:~$ jps
6323 DataNode
6597 Jps
6503 NodeManager
hduser@slave1:~$
```

On lance la commande jps dans chaque cluster

```
hduser@slave2:/usr/local/hadoop_store/hdfs$ cd ~
hduser@slave2:~$ jps
6576 NodeManager
6408 DataNode
6670 Jps
hduser@slave2:~$
```

- Vérification de l'installation : On peut maintenant accéder à l'interface web Hadoop .

Activities Firefox Web Browser 01:48

Namenode information X +

localhost:9870/dfshealth.html#tab-datanode 90% ⚡ ⚡ ⚡

Entering Maintenance In Maintenance In Maintenance & dead

### Datanode usage histogram

Disk usage of each DataNode (%)

### In operation

Show 25 entries Search:

Node	Http Address	Last contact	Last Block Report	Capacity	Blocks	Block pool used	Version
slave1:9866 (192.168.0.2:9866)	http://slave1:9864	2s	6m	9.78 GB	0	24 KB (0%)	3.2.1
slave2:9866 (192.168.0.3:9866)	http://slave2:9864	2s	6m	9.78 GB	0	24 KB (0%)	3.2.1

Showing 1 to 2 of 2 entries

Previous 1 Next 38

- On peut ajouter le noeud hadoopmaster comme datanode aussi, pour cela les fichiers suivants doivent être modifiés comme suit :

```
hduser@hadoopmaster:/usr/local/hadoop/etc/hadoop$ vim hdfs-site.xml
hduser@hadoopmaster:/usr/local/hadoop/etc/hadoop$
```

```
<!-- Put site-specific property overrides in this file. -->

<configuration>
  <property>
    <name>dfs.replication</name>
    <value>3</value>
  </property>
  <property>
    <name>dfs.namenode.name.dir</name>
    <value>file:/usr/local/hadoop_store/hdfs/namenode</value>
  </property>
  <property>
    <name>dfs.datanode.data.dir</name>
    <value>file:/usr/local/hadoop_store/hdfs/datanode</value>
  </property>
</configuration>
~
-- INSERT --
```

```
hduser@slave1:/usr/local/hadoop/etc/hadoop$ sudo vim hdfs-site.xml
[sudo] password for hduser:
hduser@slave1:/usr/local/hadoop/etc/hadoop$
```

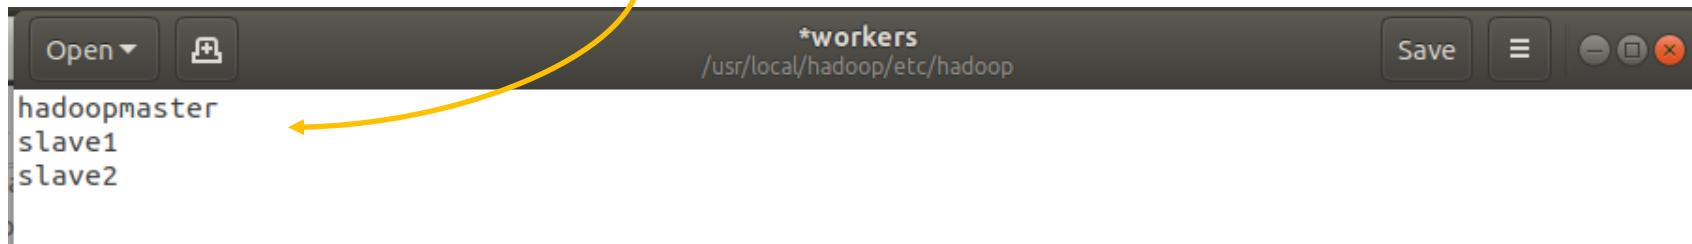
```
<configuration>
  <property>
    <name>dfs.replication</name>
    <value>3</value>
  </property>
  <property>
    <name>dfs.datanode.data.dir</name>
    <value>file:/usr/local/hadoop_store/hdfs/datanode</value>
  </property>
</configuration>
```

```
0070 3ps
hduser@slave2:~$ cd /usr/local/hadoop/etc/hadoop
hduser@slave2:/usr/local/hadoop/etc/hadoop$ vim hdfs-site.xml
hduser@slave2:/usr/local/hadoop/etc/hadoop$
```

```
<configuration>
  <property>
    <name>dfs.replication</name>
    <value>3</value>
  </property>
  <property>
    <name>dfs.datanode.data.dir</name>
    <value>file:/usr/local/hadoop_store/hdfs/datanode</value>
  </property>
</configuration>
~
~
~
~
~
" hdfs-site.xml " 28L, 971C
```

- Modification du fichier workers (slaves) dans hadoopmaster, slave1 et slave2

```
hduser@hadoopmaster:~$ cd /usr/local/hadoop/etc/hadoop
hduser@hadoopmaster:/usr/local/hadoop/etc/hadoop$ gedit workers
hduser@hadoopmaster:/usr/local/hadoop/etc/hadoop$
```



```
hduser@hadoopmaster:/usr/local/hadoop/etc/hadoop$ cd /usr/local/hadoop_store/
hduser@hadoopmaster:/usr/local/hadoop_store$ rm -rf *
hduser@hadoopmaster:/usr/local/hadoop_store$ mkdir -p /usr/local/hadoop_store/hdfs/namenode
hduser@hadoopmaster:/usr/local/hadoop_store$ mkdir -p
hduser@hadoopmaster:/usr/local/hadoop_store$ 
hduser@hadoopmaster:/usr/local/hadoop_store$ 
hduser@hadoopmaster:/usr/local/hadoop_store$ mkdir -p /usr/local/hadoop_store/hdfs/datanode
hduser@hadoopmaster:/usr/local/hadoop_store$ chown -R hduser /usr/local/hadoop_store/
hduser@hadoopmaster:/usr/local/hadoop_store$
```

- On formate le namenode.

```
hduser@hadoopmaster:/usr/local/hadoop_store$ hdfs namenode -format
2020-11-20 01:27:52,267 INFO namenode.NameNode: STARTUP_MSG:
/*****
STARTUP_MSG: Starting NameNode
STARTUP_MSG: host = hadoopmaster/192.168.0.1
STARTUP_MSG: args = [-format]
STARTUP_MSG: version = 3.2.1
```

```
2020-11-20 01:27:57,782 INFO namenode.FSNamesystem: Retry cache on namenode is enabled
2020-11-20 01:27:57,787 INFO namenode.FSNamesystem: Retry cache will use 0.03 of total heap and retry cache entry expiry time is 600000 millis
2020-11-20 01:27:57,803 INFO util.GSet: Computing capacity for map NameNodeRetryCache
2020-11-20 01:27:57,807 INFO util.GSet: VM type      = 64-bit
2020-11-20 01:27:57,808 INFO util.GSet: 0.029999999329447746% max memory 481.4 MB = 147.9 KB
2020-11-20 01:27:57,809 INFO util.GSet: capacity      = 2^14 = 16384 entries
2020-11-20 01:27:58,040 INFO namenode.FSImage: Allocated new BlockPoolId: BP-1333465830-192.168.0.1-1605835677999
2020-11-20 01:27:58,188 INFO common.Storage: Storage directory /usr/local/hadoop_store/hdfs/namenode has been successfully formatted.
2020-11-20 01:27:58,387 INFO namenode.FSImageFormatProtobuf: Saving image file /usr/local/hadoop_store/hdfs/namenode/current/fsimage.ckpt_0000
0000000000000000 using no compression
2020-11-20 01:27:58,895 INFO namenode.FSImageFormatProtobuf: Image file /usr/local/hadoop_store/hdfs/namenode/current/fsimage.ckpt_000000000000
00000000 of size 401 bytes saved in 0 seconds .
2020-11-20 01:27:58,977 INFO namenode.NNStorageRetentionManager: Going to retain 1 images with txid >= 0
2020-11-20 01:27:59,041 INFO namenode.FSImage: FSImageSaver clean checkpoint: txid=0 when meet shutdown.
2020-11-20 01:27:59,042 INFO namenode.NameNode: SHUTDOWN_MSG:
/*****
SHUTDOWN_MSG: Shutting down NameNode at hadoopmaster/192.168.0.1
*****/
```

- Démarrage de Hadoop

```
hduser@hadoopmaster:/usr/local/hadoop_store$ start-dfs.sh
Starting namenodes on [hadoopmaster]
Starting datanodes
Starting secondary namenodes [hadoopmaster]
2020-11-20 01:30:09,586 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
hduser@hadoopmaster:/usr/local/hadoop_store$ start-yarn.sh
Starting resourcemanager
Starting nodemanagers
```

- On vérifie le bon fonctionnement de tous les noeuds du cluster

```
hduser@hadoopmaster:/usr/local/hadoop/etc/hadoop$ hdfs dfsadmin -report
2020-11-20 17:54:50,905 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classe
 applicable
Configured Capacity: 20999348224 (19.56 GB)
Present Capacity: 2034827264 (1.90 GB)
DFS Remaining: 2034778112 (1.90 GB)
DFS Used: 49152 (48 KB)
DFS Used%: 0.00%
Replicated Blocks:
    Under replicated blocks: 0
    Blocks with corrupt replicas: 0
    Missing blocks: 0
    Missing blocks (with replication factor 1): 0
    Low redundancy blocks with highest priority to recover: 0
    Pending deletion blocks: 0
Erasure Coded Block Groups:
    Low redundancy block groups: 0
    Block groups with corrupt internal blocks: 0
    Missing block groups: 0
    Low redundancy blocks with highest priority to recover: 0
    Pending deletion blocks: 0
-----
Live datanodes (2):
Name: 192.168.0.2:9866 (slave1)
Hostname: slave1
Decommission Status : Normal
Configured Capacity: 10499674112 (9.78 GB)
DFS Used: 24576 (24 KB)
Non DFS Used: 8932872192 (8.32 GB)
DFS Remaining: 1013235712 (966.30 MB)
DFS Used%: 0.00%
DFS Remaining%: 9.65%
```

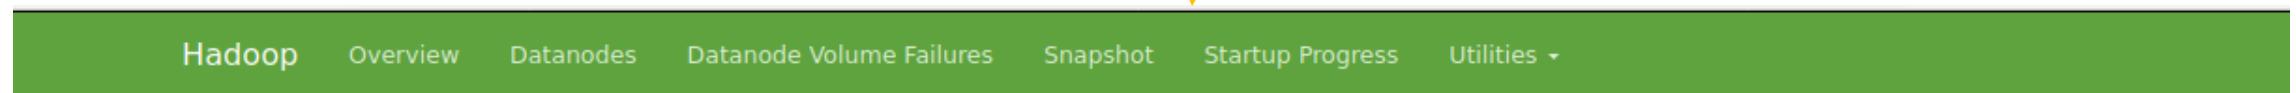
- On répète les mêmes étapes décrites dans la section (II) : Exécution d'un programme Map/Reduce dans un cluster à nœud unique.

```
hduser@hadoopmaster:/usr/local/hadoop/etc/hadoop$ cd ..
hduser@hadoopmaster:/usr/local/hadoop/etc$ cd ..
hduser@hadoopmaster:/usr/local/hadoop$ bin/dfs -put /home/hduser/Documents/code/poeme.txt /
2020-11-20 18:16:08,673 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
2020-11-20 18:16:12,889 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
```

```
hduser@hadoopmaster:/usr/local/hadoop$ cd /home/hduser/Documents/code/
hduser@hadoopmaster:~/Documents/code$ hadoop jar wcount.jar org.hadoop.wordcount.WCount /poeme.txt /results
2020-11-20 18:30:25,482 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
2020-11-20 18:30:29,293 INFO impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2020-11-20 18:30:29,645 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2020-11-20 18:30:29,645 INFO impl.MetricsSystemImpl: JobTracker metrics system started
2020-11-20 18:30:29,645 INFO input.FileInputFormat: Total input files to process : 1
```

```
    File Output Format Counters
        Bytes Written=2823
2020-11-20 18:30:38,610 INFO mapred.LocalJobRunner: Finishing task: attempt_local1575675484_0001_r_000000_0
2020-11-20 18:30:38,612 INFO mapred.LocalJobRunner: reduce task executor complete.
2020-11-20 18:30:39,577 INFO mapreduce.Job: map 100% reduce 100%
2020-11-20 18:30:39,579 INFO mapreduce.Job: Job job_local1575675484_0001 completed successfully
2020-11-20 18:30:39,706 INFO mapreduce.Job: Counters: 26
    File System Counters
        FILE: Number of bytes read=28084
        FILE: Number of bytes written=1078560
```

```
hduser@hadoopmaster:~/Documents/code$ hadoop fs -ls /results
2020-11-20 18:33:29,739 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r--  3 hduser supergroup          0 2020-11-20 18:30 /results/_SUCCESS
-rw-r--r--  3 hduser supergroup    2823 2020-11-20 18:30 /results/part-r-00000
```



## Browse Directory

	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
<input type="checkbox"/>	-rw-r--r--	hduser	supergroup	1.63 KB	Nov 20 19:16	3	128 MB	poeme.txt	
<input type="checkbox"/>	drwxr-xr-x	hduser	supergroup	0 B	Nov 20 19:30	0	0 B	results	

Showing 1 to 2 of 2 entries

Previous 1 Next

- Affichage de résultats

```
luser@hadoopmaster:~/Documents/code$ hadoop fs -cat /results/part-r-00000
2020-11-20 18:34:03,666 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
2020-11-20 18:34:07,617 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
```

a 6 occurrences.  
adoraien 1 occurrences.  
ailes 1 occurrences.  
aima 1 occurrences.  
amour 1 occurrences.  
au 11 occurrences.  
bas 1 occurrences.  
belle 1 occurrences.  
bles 1 occurrences.  
bras 1 occurrences.  
bretagne 1 occurrences.  
brula 1 occurrences.  
celle 1 occurrences.  
celui 20 occurrences.  
cette 1 occurrences.  
chancelle 1 occurrences.  
chapelle 1 occurrences.  
ciel 10 occurrences.  
citadelle 1 occurrences.  
clarte 1 occurrences.  
coeur 2 occurrences.  
combat 1 occurrences.  
comment 1 occurrences.  
commun 1 occurrences.

nouvelle 1 occurrences.  
ny 10 occurrences.  
ou 3 occurrences.  
par 1 occurrences.  
pas 11 occurrences.  
passent 1 occurrences.  
plus 2 occurrences.  
pour 1 occurrences.  
prefere 1 occurrences.  
prison 1 occurrences.  
prisonniere 1 occurrences.  
qua 1 occurrences.  
quand 2 occurrences.  
quaucun 1 occurrences.  
que 2 occurrences.  
quelle 1 occurrences.  
querelles 1 occurrences.  
qui 25 occurrences.  
quil 1 occurrences.  
quimporte 1 occurrences.  
raisin 1 occurrences.  
rats 1 occurrences.  
rebelle 2 occurrences.  
rechantera 1 occurrences.  
repetant 1 occurrences.  
reseda 1 occurrences.  
rose 1 occurrences.  
rouge 1 occurrences.  
ruisselle 1 occurrences.  
saison 1 occurrences.  
sang 1 occurrences.  
sanglots 1 occurrences.  
sappelle 1 occurrences.  
se 1 occurrences.

prisonniere 1 occurrences.  
qua 1 occurrences.  
quand 2 occurrences.  
quaucun 1 occurrences.  
que 2 occurrences.  
quelle 1 occurrences.  
querelles 1 occurrences.  
qui 25 occurrences.  
quil 1 occurrences.  
quimporte 1 occurrences.  
raisin 1 occurrences.  
rats 1 occurrences.  
rebelle 2 occurrences.  
rechantera 1 occurrences.  
repetant 1 occurrences.  
reseda 1 occurrences.  
rose 1 occurrences.  
rouge 1 occurrences.  
ruisselle 1 occurrences.  
saison 1 occurrences.  
sang 1 occurrences.  
sanglots 1 occurrences.  
sappelle 1 occurrences.  
se 1 occurrences.

....

**Hadoop** en tant que service (HDaaS) facilite l'approche des projets big data. Voici quelques-uns des principaux fournisseurs de services cloud Hadoop qui rendent cette course encore plus intéressante :

**1)Amazon Web Service EMR (AWS EMR)**

**2)GOOGLE CLOUD**

**3)Cloudera**

**4)Microsoft Azure- HDInsight**

**5) IBM BigInsight**

**6) MapR**

**7)CSC**



### 1) Amazon Web Service EMR (AWS EMR):

- Amazon EMR (Amazon Elastic Map Reduce) est actuellement l'un des principaux fournisseurs de services cloud Hadoop. En outre, Amazon EMR n'est pas seulement limité à Hadoop, mais fournit également des services à Spark et d'autres solutions Big Data.
- Amazon EMR facilite la création et la gestion de clusters élastiques et entièrement configurés d'instances Amazon EC2 exécutant Hadoop et d'autres applications dans l'écosystème Hadoop.

- **Avantages d' EMR :**

### **Vitesse et souplesse accrues**

Vous pouvez initialiser un cluster Hadoop de façon dynamique et rapide, ou ajouter des serveurs à votre cluster Amazon EMR existant, pour mettre les ressources plus rapidement à la disposition de vos utilisateurs et spécialistes des données. En utilisant Hadoop sur la plateforme AWS, votre organisation peut considérablement gagner en souplesse en diminuant le coût et le temps nécessaires à l'allocation des ressources à des fins d'expérimentation et de développement.

### **Réduction de la complexité administrative**

La configuration d'Hadoop, la mise en réseau, l'installation du serveur, la configuration de la sécurité et la maintenance administrative courante peuvent s'avérer des tâches complexes et difficiles. Amazon EMR étant un service géré, la solution répond aux exigences de votre infrastructure Hadoop pour vous permettre de vous concentrer sur votre activité principale.

## **Vous pouvez exploiter facilement d'autres services AWS.**

Vous pouvez facilement intégrer votre environnement Hadoop à d'autres services tels qu'Amazon S3, Amazon Kinesis, Amazon Redshift et Amazon DynamoDB pour permettre le transfert de données, les workflows et les analyses dans les nombreux services de la plate-forme AWS. En outre, vous pouvez utiliser le catalogue de données AWS Glue comme référentiel de métadonnées géré pour Apache Hive et Apache Spark.

## **Payez pour les clusters uniquement quand vous en avez besoin**

De nombreuses tâches Hadoop sont par nature irrégulières. Par exemple, une tâche ETL peut être exécutée une fois par heure, par jour ou par mois, tandis que des tâches de modélisation effectuées pour des sociétés financières ou des séquençages génétiques peuvent avoir lieu seulement quelques fois par an. En utilisant Hadoop sur Amazon EMR, vous pouvez facilement lancer ces clusters de charges de travail, enregistrer les résultats et supprimer vos ressources Hadoop lorsque vous n'en avez plus besoin, pour éviter les coûts d'infrastructure inutiles.

- Les inconvénients :

### **Absence of Hadoop Management Console**

AWS ne fournit aucune console de gestion similaire à Ambari ou Cloudera Manager d'Apache, pour EMR. Il est donc difficile de gérer et de surveiller divers services Hadoop sur un cluster en cours d'exécution.

### **No High Availability for Master Node**

Le nœud maître du cluster EMR n'est pas configuré pour la haute disponibilité, ce qui en fait le point unique d'échec.

### **Automatic Replacement of Unhealthy Nodes**

Le service EMR surveille en permanence les nœuds esclaves et remplace de manière transparente tout nœud malsain. Tout en offrant un avantage de maintenance distinct aux administrateurs, cette fonctionnalité risque également la perte de données résidant sur le nœud malsain.



## 2) GOOGLE CLOUD

Grâce à Dataproc, vous pouvez créer une ou plusieurs instances Compute Engine pouvant se connecter à une instance de Cloud Bigtable et exécuter des tâches Hadoop. On utiliser Dataproc pour automatiser les tâches suivantes :

- Installer Hadoop et le client HBase pour Java
- Configurer Hadoop et Cloud Bigtable
- Définir les champs d'application d'autorisations appropriés pour Cloud Bigtable

Après avoir créé un cluster Dataproc, vous pouvez l'utiliser pour exécuter des tâches Hadoop qui lisent et écrivent des données depuis et vers Cloud Bigtable.

### 3) Cloudera



Cloudera est également l'une des rares entreprises qui fournissent la mise en place complète pour Hadoop. En fait, Cloudera est le plus célèbre de tous.

CDH est un logiciel Apache 100% open source et est la seule solution Hadoop à offrir un traitement par lots uniifié, sql interactif, recherche interactive, et des contrôles d'accès basés sur les rôles.

Vous pouvez commencer par Cloudera Free et utiliser les informations d'identification par défaut de Cloudera pour commencer. Si votre système dispose de 10 Go de RAM, vous pouvez également faire l'expérience de Cloudera Manager.



#### 4) Microsoft Azure- HDInsight :

HDInsight est la distribution Hadoop alimentée par le cloud. Il a été conçu pour mettre à l'échelle et traiter les données à partir de téraoctets à pétaoctets.

Il s'agit d'une offre Hadoop cloud entièrement gérée qui fournit des clusters analytiques open source optimisée pour Spark, Hive, MapReduce, HBase, Storm, Kafka et R Server avec 99,9% SLA. Toutes ces technologies big data et applications ISV sont facilement déployables en tant que clusters gérés avec sécurité et surveillance au niveau de l'entreprise.

Les serveurs sont facilement configurables avec de nombreux outils de productivité tels que Datameer, Cask, AtScale et StreamSets.

Le service cloud Hadoop de Microsoft Azure est facile à gérer et à gérer pour les administrateurs. Avec HDInsight, vous pouvez traiter des données non structurées ou semi-structurées à partir de clics web, de médias sociaux, de journaux de serveurs, d'appareils et de capteurs, et plus encore.

## 5) IBM BigInsight:



IBM BigInsight est un important fournisseur de services cloud Hadoop qui fournit le service cloud sur l'infrastructure cloud mondiale IBM SoftLayer.

IBM Infosphere BigInsight ne nécessite aucune infrastructure locale et il prend en charge Big SQL, Big Sheets, et l'analyse de texte et plus encore, IBM affirme.

Voici quelques-unes des caractéristiques de l'édition standard IBM BigInsight-

- Entièrement intégré, entièrement compatible – Installation intégrée d'Apache Hadoop et des composants open source associés de l'écosystème Apache Hadoop qui est testé et préconfiguré.
- Comprend Jaql, un langage de requête déclaratif, pour faciliter l'analyse des données structurées et non structurées.
- Fournit une console de gestion web pour faciliter l'administration et les vues en temps réel.
- Comprend BigSheets, un outil d'analyse et de visualisation web doté d'une interface familière, semblable à une feuille de calcul, qui permet d'analyser facilement de grandes quantités de données et des tâches de collecte de données à long terme.
- Inclut Big SQL, un moteur de requête SQL natif qui permet l'accès SQL aux données stockées dans BigInsights, en tirant parti de MapReduce pour des jeux de données complexes et un accès direct pour les requêtes plus petites.

## 6) MapR :



MapR fournit la distribution complète pour Hadoop et est un environnement complet travaillant sur Hadoop 2.0. MapR est également l'un des plus grands fournisseurs d'Apache Hadoop.

La distribution complète de MapR comprend- Apache Hive, Apache Pig, Cascading, Apache HCatalog, Apache HBase™, Apache Oozie, Apache Flume, Apache Sqoop, Apache Mahout et Apache Whirr.

Voici quelques-unes des caractéristiques de MapR:

- Rentabilité
- tolérant aux défauts
- flexible
- évolutif

## 7) CSC :

CSC est également l'un des principaux fournisseurs de services cloud Big Data Hadoop dans le monde. Ils fournissent un programme intégré entièrement géré.

CSC Big Data Platform as a Service (BDPaaS) aide les entreprises à franchir ces obstacles et à tirer de la valeur de leurs données beaucoup plus rapidement. Avec BDPaaS, les entreprises peuvent rapidement développer, sécuriser et déployer des applications big data et analytiques de nouvelle génération avec une plate-forme centralisée basée sur l'abonnement qui utilise des outils d'analyse, une infrastructure et des logiciels de pointe.

Le logiciel comprend une prise en charge intégrée de Cloudera, Hortonworks, DataStax, Spark, Pentaho, Qlik, Tableau, R, Python, et plus encore.

- Voici quelques-unes des caractéristiques de BDPaaS-
- Augmenter les taux de réussite
- Accélérer l'heure de la valeur
- Réduire les coûts grâce à des applications open-source
- Protéger les données grâce à la sécurité d'entreprise multicouches
- Activer le développement rapide des applications big data et bien d'autres

Fin

Merci